

# Digital Image Processing

Berlin University of Technology (TUB),  
Computer Vision and Remote Sensing Group  
Berlin, Germany



# Digital Image Processing

## (One) Goal of Computer Vision:

Automatic Understanding of digital Images!

Image is distorted?



→ Image restoration (e.g. Wiener filter)

Image has still bad quality?




→ Image enhancement (e.g. Histogram equalization)

How to describe an image (or the image content)?

→ Just pixel intensity?!

# Image Features

**Digital Image  
Processing**



**Automatic Image  
Analysis**

Image → Image

[DIP Winter term]

How to get a meaningful  
image description?

Image → Features

Features → “Understanding”

[AIA Summer term]  
[PCV Winter term]

How to use this image  
description to infer  
information about image  
content?

# Image Features

Color




Intensity






# Image Features

Edges



Segments



# Image Features

Texture





Interest Points








# Interest Points



# Interest Points




K\_Y P\_I N T ?





# Interest Points




K\_Y P\_I N T ?

Keypoint Detector:

- clear mathematical definition
- well-defined position
- rich local information contents
- stable under perturbations
- reliable

# Interest Points



K\_Y P\_I N T ?

Keypoint Detectors:

→ Harris-Stephens

→ Förstner

→ Shi-Tomasi

→ SUSAN

→ FAST

→ SIFT

→ SURF

→ ....

# Basics

- Directional Gradients
- Covariance Matrices




# Directional Gradients

- Often only gradient magnitude is computed:
  - Use e.g. a radially symmetric filter
  - No information concerning the direction of gradients
- Now: **Directional gradients**
  - Convolution with suitable filters, e.g.  $G_x$  and  $G_y$
  - **Image**  $\otimes$   **$G_x$**  → Gradient in x direction
  - **Image**  $\otimes$   **$G_y$**  → Gradient in y direction
- Each pixel is associated with a gradient vector  **$\mathbf{g} = (g_x, g_y)^T$**



# Directional Gradients



- Gradient magnitude:  $|g| = \sqrt{g_x^2 + g_y^2}$

- Gradient direction:  $\theta = \tan^{-1} \left( \frac{g_y}{g_x} \right)$

→ Direction in which intensity increases quickest

# Directional Gradients

- Commonly Used: Composition of differential operator and low-pass
- E.g. derivatives of the normal distribution:

$$G(x, y; \sigma) = \frac{1}{2\pi\sigma^2} \exp\left(-\frac{x^2 + y^2}{2\sigma^2}\right)$$


$$G_x(x, y) = \frac{\partial}{\partial x} G(x, y; \sigma) = \frac{-x}{2\pi\sigma^4} \exp\left(-\frac{x^2 + y^2}{2\sigma^2}\right) = \frac{-x}{\sigma^2} G(x, y; \sigma)$$

$$G_y(x, y) = \frac{\partial}{\partial y} G(x, y; \sigma) = \frac{-y}{2\pi\sigma^4} \exp\left(-\frac{x^2 + y^2}{2\sigma^2}\right) = \frac{-y}{\sigma^2} G(x, y; \sigma)$$


- $\sigma$ : Scale and noise sensitivity
  - $\sigma$  small: Small structures discernible, noise/texture preserved
  - $\sigma$  large: Large structures emphasized, noise suppressed



# Directional Gradients




# Directional Gradients



$g_x$



$g_y$



$|g|$

# Directional Gradients



$g_x$



$g_y$



$|g|$




# Basics

- Directional Gradients
- Covariance Matrices

# Covariance Matrices

- Variance of scalars  $\{x_1, x_2, x_3, \dots, x_N\}$ : Measures dispersion around mean  $\mu$



# Covariance Matrices

- For sets of vectors  $\{\mathbf{x}_1, \mathbf{x}_2, \mathbf{x}_3, \dots, \mathbf{x}_N\}$  ( $\mathbf{x}_j$  M-dimensional):

- Mean 
$$\mu = \frac{1}{N} \sum_{j=1}^N \mathbf{x}_j$$

- Covariance 
$$\Sigma = \frac{1}{N} \sum_{j=1}^N (\mathbf{x}_j - \mu)(\mathbf{x}_j - \mu)^T$$


- For  $\mathbf{x}_j = (x_{j,1}, x_{j,2}, x_{j,3}, \dots, x_{j,M})^T$ :

$$\Sigma = \frac{1}{N} \begin{bmatrix} \sum_{j=1}^N (x_{j,1} - \mu_1)^2 & \sum_{j=1}^N (x_{j,1} - \mu_1)(x_{j,2} - \mu_2) & \cdots & \sum_{j=1}^N (x_{j,1} - \mu_1)(x_{j,M} - \mu_M) \\ \sum_{j=1}^N (x_{j,2} - \mu_2)(x_{j,1} - \mu_1) & \sum_{j=1}^N (x_{j,2} - \mu_2)^2 & \cdots & \sum_{j=1}^N (x_{j,2} - \mu_2)(x_{j,M} - \mu_M) \\ \vdots & \vdots & \ddots & \vdots \end{bmatrix}$$

- Diagonal: Variance along individual dimensions
- Otherwise: Correlation between dimensions




# Covariance Matrices



$$\Sigma = \begin{pmatrix} 0.9976 & -0.0187 \\ -0.0187 & 0.9700 \end{pmatrix}$$

Both dimensions almost uncorrelated




$$\Sigma = \begin{pmatrix} 0.4998 & 0.4625 \\ 0.4625 & 0.5054 \end{pmatrix}$$

Both dimensions strongly correlated

# Covariance Matrices

- Component of a vector  $\mathbf{x}$  in direction  $\mathbf{r}$ :



- $\mathbf{r} = (\cos \theta, \sin \theta)^T$
- $x_r = \mathbf{r}^T \mathbf{x}$
- $x_r$ : Scalar, component of  $\mathbf{x}$  along  $\mathbf{r}$

- Mean and variance of vectors  $\{\mathbf{x}_1, \mathbf{x}_2, \mathbf{x}_3, \dots, \mathbf{x}_N\}$  along direction  $\mathbf{r}$ :

$$\mu_r = \frac{1}{N} \sum_j x_{r,j} = \frac{1}{N} \sum_j \mathbf{r}^T \mathbf{x}_j = \mathbf{r}^T \frac{1}{N} \sum_j \mathbf{x}_j = \mathbf{r}^T \boldsymbol{\mu}$$

$$\begin{aligned} \Sigma_r &= \frac{1}{N} \sum_j (x_{r,j} - \mu_r)(x_{r,j} - \mu_r)^T = \frac{1}{N} \sum_j (\mathbf{r}^T \mathbf{x}_j - \mathbf{r}^T \boldsymbol{\mu})(\mathbf{r}^T \mathbf{x}_j - \mathbf{r}^T \boldsymbol{\mu})^T \\ &= \frac{1}{N} \sum_j \mathbf{r}^T (\mathbf{x}_j - \boldsymbol{\mu})(\mathbf{x}_j - \boldsymbol{\mu})^T \mathbf{r} = \mathbf{r}^T \Sigma \mathbf{r} \end{aligned}$$

- The covariance matrix determines the variance in all directions!

# Covariance Matrices

Task: Find directions with maximal variance, i.e.:  $r^T \Sigma r = \max$

Solution:

$$\Sigma = \mathbf{V} \mathbf{D} \mathbf{V}^T$$


$\begin{pmatrix} \uparrow & \uparrow & \dots \\ \mathbf{v}_1 & \mathbf{v}_2 & \dots \\ \downarrow & \downarrow & \dots \end{pmatrix}$ 

Eigenvectors

$\begin{pmatrix} l_1 & 0 & \dots \\ 0 & l_2 & \dots \\ \vdots & \vdots & \ddots \end{pmatrix}$ 

Eigenvalues (diagonal)

- In the eigenbasis, dimensions of vectors  $\mathbf{x}$  are not correlated



- Variance along in the directions  $\mathbf{v}_1, \mathbf{v}_2 \dots : l_1, l_2, \dots$
- Standard deviations in other directions form an ellipse with major/minor axes along eigenvectors with deviations given by eigenvalues


# Structure Tensor

- For each pixel, the structure tensor **A** is defined as:

$$\mathbf{A} = \sum_W \mathbf{g} \mathbf{g}^T = \begin{pmatrix} \sum_W g_x^2 & \sum_W g_x g_y \\ \sum_W g_y g_x & \sum_W g_y^2 \end{pmatrix}$$

- $W$  denotes the neighborhood of the pixel considered
  - In this exercise: Gaussian window with std-dev  $N_w$  around the pixel
- A** is a covariance matrix computed assuming  $\mu = 0$ 
  - **A** describes the distribution of gradients around  $\mathbf{g} = (0,0)^T$
- For covariance  $\Sigma$ :  $\mathbf{r}^T \Sigma \mathbf{r} = \text{Variance along } \mathbf{r}$
- For the structure tensor **A**:  $\mathbf{r}^T \mathbf{A} \mathbf{r} = (\text{Squared})$  gradient magnitude along  $\mathbf{r}$


# Structure Tensor



- $\mathbf{v}_1$ : Direction with the greatest gradient magnitude (max. eigenvalue  $l_1$ )
  - Gradient direction  $\mathbf{v}_1$  dominates neighborhood  $W$
- $l_1$ : Total (squared) gradient magnitude along direction  $\mathbf{v}_1$
- $\mathbf{v}_2$ : Direction with the smallest gradient magnitude
  - Gradient direction  $\mathbf{v}_2$  is rare in neighbourhood  $W$
- Gradient magnitude as a function of direction describes an ellipse with major/minor axes along  $\mathbf{v}_1$  und  $\mathbf{v}_2$



# Structure Tensor



$$\mathbf{A}^{-1} = (\mathbf{V} \mathbf{D} \mathbf{V}^T)^{-1} = \mathbf{V} \mathbf{D}^{-1} \mathbf{V}^T$$

$$\mathbf{D}^{-1} = \begin{pmatrix} \frac{1}{l_2} & 0 \\ 0 & \frac{1}{l_1} \end{pmatrix} = \begin{pmatrix} \lambda_1 & 0 \\ 0 & \lambda_2 \end{pmatrix}$$

- Eigenvectors remain unchanged
- Eigenvalues are inverted
- Small eigenvalues  $\lambda_1$  und  $\lambda_2$  indicate strong gradients in the neighborhood
- If  $\lambda_1$  and  $\lambda_2$  are large, the image is homogeneous

# Förstner Operator

- The structure tensor can be used to derive salient information:
- Weight **w**: Strength of gradients in the neighborhood

$$w = \frac{1}{\text{tr}(\mathbf{A}^{-1})} = \frac{1}{\lambda_1 + \lambda_2} = \frac{\det(\mathbf{A})}{\text{tr}(\mathbf{A})} \quad w > 0$$


- w large:  $\lambda_1$  und  $\lambda_2$  small, i.e. strong gradients in the neighborhood
- $w_{\min} = 0.5, \dots, 1.5 \cdot \bar{w}$ ,  $\bar{w}$  is the mean of w over whole image

- Isotropy **q**: Measures the uniformity of gradient directions in the neighbourhood

$$q = 1 - \left( \frac{\lambda_1 - \lambda_2}{\lambda_1 + \lambda_2} \right)^2 = \frac{4\det(\mathbf{A})}{\text{tr}(\mathbf{A})^2} \quad 0 \leq q \leq 1$$

- **q** small: Gradients occur primarily in one direction
- $q_{\min} = 0.5, \dots, 0.75$



# Förstner Operator



Original image

# Förstner Operator



1. Gradient in x-direction



Gradient in x-direction

# Förstner Operator

1. Gradient in x- and y-direction





Gradient in y-direction



# Förstner Operator


1. Gradient in x- and y-direction
2.  $g_x \cdot g_x$



$$g_x \cdot g_x$$

# Förstner Operator

1. Gradient in x- and y-direction
2.  $g_x \cdot g_x, g_y \cdot g_y$





$$g_y \cdot g_y$$




# Förstner Operator

1. Gradient in x- and y-direction
2.  $g_x \cdot g_x, g_y \cdot g_y, g_x \cdot g_y$



$$g_x \cdot g_y$$

# Förstner Operator


1. Gradient in x- and y direction 
2.  $g_x \cdot g_x, g_y \cdot g_y, g_x \cdot g_y$
3. Average (Gaussian Window)



Averaged (smoothed)  $g_x \cdot g_y$



# Förstner Operator



1. Gradient in x- and y direction
2.  $g_x \cdot g_x, g_y \cdot g_y, g_x \cdot g_y$
3. Average (Gaussian Window)
4. Trace of structure tensor

tr(A)





# Förstner Operator

1. Gradient in x- and y direction
2.  $g_x \cdot g_x, g_y \cdot g_y, g_x \cdot g_y$
3. Average (Gaussian Window)
4. Trace of structure tensor
5. Determinant of structure tensor



|A|



# Förstner Operator

1. Gradient in x- and y direction
2.  $g_x \cdot g_x, g_y \cdot g_y, g_x \cdot g_y$
3. Average (Gaussian Window)
4. Trace of structure tensor
5. Determinant of structure tensor
6. weight calculation





Weight  $w$




# Förstner Operator

1. Gradient in x- and y direction
2.  $g_x \cdot g_x, g_y \cdot g_y, g_x \cdot g_y$
3. Average (Gaussian Window)
4. Trace of structure tensor
5. Determinant of structure tensor
6. weight calculation
7. isotropy calculation



Isotropy q

# Förstner Operator



1. Gradient in x- and y direction
2.  $g_x \cdot g_x, g_y \cdot g_y, g_x \cdot g_y$
3. Average (Gaussian Window)
4. Trace of structure tensor
5. Determinant of structure tensor
6. weight calculation
7. isotropy calculation
8. Keypoint extration
  - weight > weight threshold
  - isotropy > isotropy threshold
  - weight is local maximum

Keypoints



# 5. Exercise – Given

```
int main(int argc, char** argv)
```

- Loads image, extracts and shows keypoints
- argv[1] == path to image
- argv[2] == scale of kernel (std-dev) for directional gradients
- argv[3] == scale of kernel (std-dev) for neighborhood

```
unsigned getOddKernelSizeForSigma(float sigma)
```

sigma    Standard deviation

return   Kernel size to use (always odd)

- Institutionally mandated "correct" kernel size
- Makes unit testing easier for me

```
bool isLocalMaximum(const cv::Mat_<float>& img, int x, int y)
```

img            input image

x,y            pixel location. Note: x == col, y == row

return        true if value at (x,y) in img is locally maximal

- Checks if all neighbors are smaller



# 5. Exercise – ToDo

```
cv::Mat_<float> createGaussianKernel1D(float sigma)
```



sigma          std-dev of filter kernel  
return          1D gaussian kernel (horizontal layout)

- Computes 1D Gaussian kernel for separable convolutions
- Compute kernel size using `getOddKernelSizeForSigma`
- Copy/Adapt from previous homework

```
Mat separableFilter(Mat& src, Mat& kernelX, Mat& kernelY)
```



src            Image to filter  
kernelX       1D kernel to apply horizontally (kernel in horizontal layout)  
kernelY       1D kernel to apply vertically (kernel in horizontal layout)  
return        Filtered image (same size)

- Computes separable convolution
- Note that different kernels can be used for horizontal and vertical passes
- Copy/Adapt from previous homework

# 5. Exercise – ToDo

```
cv::Mat_<float> createFstDevKernel1D(float sigma)
```



sigma          std-dev of filter kernel (first derivative of Gaussian)  
return          the created kernel

- Generates kernel that corresponds to the first derivative of a Gaussian

$$G_x(x) = \frac{\partial}{\partial x} G(x; \sigma) = \frac{-x}{2\pi\sigma^4} \exp\left(-\frac{x^2}{2\sigma^2}\right)$$

```
void calculateDirectionalGradients(cv::Mat& img, float sigmaGrad,  
                                   cv::Mat_<float>& gradX, cv::Mat_<float>& gradY)
```

img            input image  
sigmaGrad    std-dev of Gaussians  
gradX        Output matrix for x-components of per pixel gradients  
gradY        Output matrix for y-components of per pixel gradients



- Computes directional gradients via **separable** convolution
- For x-component: Convolve horizontally with **derivative of Gaussian** and vertically with **normal Gaussian**
- For y-component: The other way around


# 5. Exercise – ToDo

```
void calculateStructureTensor(Mat& gradX, Mat& gradY, float sigma,  
                             Mat& A00, Mat& A01, Mat& A11)
```

gradX, gradY

Input directional gradients


sigma

Std-dev for the Gaussian blur to compute the neighborhood sum. 

A00, A01, A11

Output per pixel structure tensor matrix

- Computes the structure Tensor for each pixel
- Neighborhood summation through convolution with Gaussian kernel
- Output tensor matrix elements as separate matrices

$$\mathbf{A} = \sum_W \mathbf{g} \mathbf{g}^T = \begin{pmatrix} \sum_W g_x^2 & \sum_W g_x g_y \\ \sum_W g_y g_x & \sum_W g_y^2 \end{pmatrix}$$


Use Gaussian blur instead of plain summation

# 5. Exercise – ToDo

```
void calculateFoerstnerWeightIsotropy(Mat& A00, Mat& A01, Mat& A11,  
                                     Mat& weight, Mat& isotropy)
```

A00, A01, A11	Input per pixel structure tensor matrices
weight	Output per pixel "Förstner weight"
isotropy	Output per pixel "Förstner isotropy"



- Computes per pixel the weight and isotropy
- Prevent division by zero:
  - [...] / std::max(trace, 1e-8f)
  - [...] / std::max(trace \* trace, 1e-8f)

$$w = \frac{1}{\text{tr}(\mathbf{A}^{-1})} = \frac{1}{\lambda_1 + \lambda_2} = \frac{\det(\mathbf{A})}{\text{tr}(\mathbf{A})} \quad w > 0$$

$$q = 1 - \left( \frac{\lambda_1 - \lambda_2}{\lambda_1 + \lambda_2} \right)^2 = \frac{4\det(\mathbf{A})}{\text{tr}(\mathbf{A})^2} \quad 0 \leq q \leq 1$$

# 5. Exercise – ToDo

```
vector<Vec2i> getFoerstnerInterestPoints(Mat& img,  
                                         float sigmaGrad, float sigmaNeighborhood,  
                                         float minWeight, float minIsotropy)
```

img	input image
sigmaGrad	std-dev of filter kernels for directional gradients
sigmaNeighborhood	std-dev of filter kernel for neighborhood summation
minWeight	Minimum weight of interest points as fraction of average weights
minIsotropy	Minimum isotropy of interest points
return	found keypoint locations ( <b>column, row</b> )

- Computes directional gradients, structure tensors, weights and isotropies
- Extracts pixel locations where:
  - weight is larger than computed weight threshold
  - isotropy is larger than minIsotropy
  - weight is local maximum
- Use isLocalMaximum(...) to check if weight is local maximum

$w_{min} = 0.5, \dots, 1.5 \cdot \bar{w}$ ,  $\bar{w}$  is the mean of  $w$  over whole image

# 6. Exercise – Testcode

main.cpp contains a piece of test code for exercise 6:

```
#define TEST_FOR_DIP6

#ifdef TEST_FOR_DIP6
#include <immintrin.h>
void testForDip6() {
    .....
}
#else
void testForDip6() {
}
#endif
```

This checks if stuff that we will need for exercise 6 works on your PC. If this fails during compilation or runtime, comment out "#define TEST\_FOR\_DIP6" and write me an email.

# Next Meeting

**Deadline: 21.01.2020**

**Next Meeting: 21.01.2019**

No theory questions but do remember to include input/output images in submission.